

This is episode 1 in series of blog posts that seek to explain how mental representations became a powerful tool in psychology and cognitive neuroscience from the early 20th century to the present day. This episode focusses on the dire straits the field found itself in the early 20th century and the novel quantitative paradigm that allowed the field to objectively quantify psychological phenomena (perhaps for the first time!).

The Crisis of Confidence

In the early 20th century, psychology was facing a crisis of confidence. As Joel Michell pointed out in his very well written 1999 book, Measurement in Psychology [1], psychology faced an immense pressure to be quantitative like Physics was at that time was. Note that Einstein, Bohr and Planck were all revolutionizing the way we think about the physical world at that time. The main way they were doing this by coming with new theories but with objective ways of to the phenomena. As Michell points out in his book, psychology as a field was not viewed with much respect by the scientific community felt a critical need to be more quantitative to earn legitimacy.

This is notwithstanding the fact that later in the century, the Ferguson report [2], commissioned by the British Association for the Advancement of Science released in 1940, concluded that there is no real way to perform psychological measurement thereby eviscerating the very foundation of psychology as a science - at least at the time.

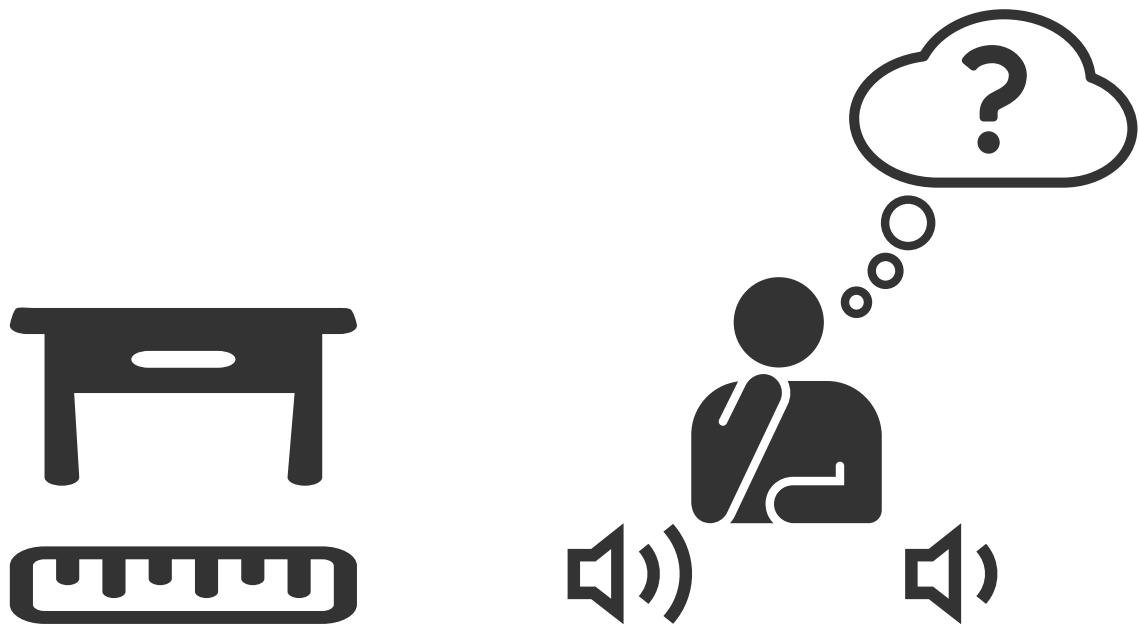
Why is similarity needed? And why not absolute values?

The main reason for this evisceration was largely due to the fact that psychology did not have a ground truth to measure against. The canonical example of “ground truth” from physics is that of length. There exists an absolute length of a ruler and of the standard measurement units (meters) for length. However, in psychology, there is no rule or scale to measure the perception of loudness of a sound perceived by an individual even if there are objective measures of loudness (decibels).

This becomes even harder when you consider concepts that do not have a physical correlate like sound does. For example, aesthetic appeal of a painting. How do you measure the beauty of a painting using an absolute scale? You cannot - at least not in a way that is objective and can be agreed upon by everyone.

In the midst of this measurement crisis for psychology, L.L. Thurstone in his 1927 paper, A Law of Comparative Judgement [3] proposed a new way to perform psychological measurement. He argued that while we as humans may not have an absolute scale, we possess a highly sensitive comparison mechanism to compare any two stimuli via a relative scale.

That is, we are not very good at saying “This sound is at 400 dB” but we are very good at saying “Sound A is louder than Sound B”. The idea that humans are more consistent in their relative judgements than absolute judgements has been replicated repeatedly since then [4].



Length is an absolute objective scale

Perception lacks a universal absolute scale

Figure 1: Absolute vs Relative Scales. Human perception does not have an absolute scale to measure against like physical concepts like length or weight. However, humans possess a high sensitive comparison mechanism to compare any two stimuli via a relative scale.

Scales: The 1D Era

While relative judgements allow one to get around the fact that there exists no absolute scale to measure against, it still doesn't allow one to account for the fact that humans are highly inconsistent in comparing stimuli due to factors like attention, fatigue etc.

The way that Thurstone was able to explain human inconsistency is via assuming that internal noise could be modelled as a Gaussian Distribution and that each time an individual made a judgement about a particular stimuli, that judgement was a draw from that distribution. More formally, if two sound stimuli μ_A and μ_B are presented to an individual where the perceived loudness is sampled from $N(\mu_A, \sigma^2)$ and $N(\mu_B, \sigma^2)$ independently where x_A and x_B are the samples from the respective distributions.

This way, each time the individual is presented with the same two stimuli, the individual essentially draws from each of the two Gaussian distributions to compare the resulting samples.

To model if the individual perceives stimulus A as louder than stimulus B , we can compute the difference between the two draws, $\delta = x_A - x_B$.

- If $\delta > 0$, the individual perceives stimulus A as louder than stimulus B .
- If $\delta < 0$, the individual perceives stimulus B as louder than stimulus A .

Consider a thought experiment with two sets of stimuli. The first comparison is between a whisper and a bomb blast where the human assesses on all trials that the bomb blast is louder than the whisper. The second is between a bomb blast and a gun firing where the human assesses that the bomb blast is louder than the gun firing only about 80% of the time. What this essentially tells you is that the loudness of the bomb blast is very far from the loudness of the whisper, while the loudness of the bomb blast is closer to the loudness of the gun firing.

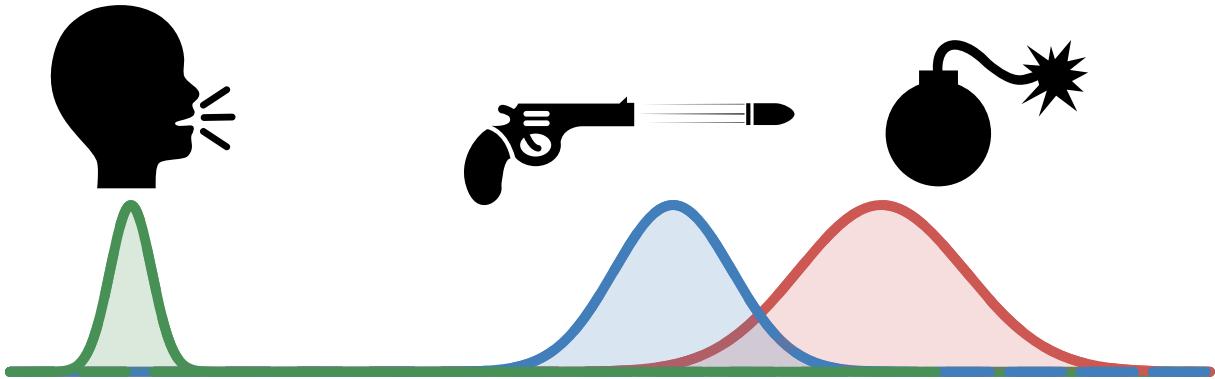


Figure 2: A whisper and a bomb blast are perceived as very different in loudness, while a bomb blast and a gun firing are perceived as less different in loudness due to increase in confusion between the stimuli.

To formalize this, note that the difference between gaussian distributions is also gaussian [5]. Therefore, the percentage of times that the individual perceives stimulus A as different from stimulus B is related to the distance between the two stimuli.

For example, if $P(A > B) = 0.5$, then the distance between the means μ_A and μ_B is 0. If $P(A > B) = 0.84$, then the distance between the means is 1 standard deviation. If $P(A > B) = 0.977$, then the distance between the means is 2 standard deviations.

If these numbers seem familiar, it is because they are essentially the same numbers that you would see in a standard normal distribution table. Basically, the Inverse Cumulative Distribution Function (CDF) or the Z-score function, $Z(\cdot)$, evaluated on the probability of A being greater than B gives you the distance between the means in terms of standard deviations. Essentially you assume that $D_{AB} = Z(P(A > B))$.

With this technique, Thurstone essentially was able to derive a quantitative scale for measuring the difference between two stimuli based on the relative judgements. He turned the natural human error in judgement into a feature of the measurement scale. (His paper additionally described formulae for cases where the standard deviation of the internal noise was not constant across stimuli. However, the same idea holds.)

But wait, we just have a collection of pairwise distances? How do we go from that to a scale? You might be thinking of Multidimensional Scaling (MDS) here. But that was created only in late 1950s and early 1960s so Thurstone couldn't have used it (We will go in detail about MDS in a future episode). And also, this is only a One Dimensional Scale so MDS is overkill when something far simpler would suffice.

The algorithm is actually pretty simple. They created a pairwise distance matrix \mathbf{D} where D_{ij} is the distance (Z-score) between stimulus i and stimulus j after aggregating over all of the data. Then, to find the scale they just averaged over the columns of \mathbf{D} to get a single value for each stimulus. This is akin to taking a "centroid" of the stimulus by comparing its relative score and distance to all other stimuli. By averaging how much 'better' or 'worse' a crime is compared to every other crime in the set, we find its unique coordinate on the moral map. Then, if you sort the stimuli based on this centroid, you get the scale.

$$D = \begin{pmatrix} Z(P(0,0)) & \dots & Z(P(0,n-1)) \\ \vdots & \ddots & \vdots \\ Z(P(n-1,0)) & \dots & Z(P(n-1,n-1)) \end{pmatrix} = \begin{pmatrix} D_{0,0} & \dots & D_{0,n-1} \\ \vdots & \ddots & \vdots \\ D_{n-1,0} & \dots & D_{n-1,n-1} \end{pmatrix} \xrightarrow{\text{averaging}} \xrightarrow{\text{Centroid Calculation}} (s_0 \ s_1 \ \dots \ s_{n-1})$$

Figure 3: Making the distance matrix from the the Z scored pairwise probabilities and then averaging over the columns to get the scale for each stimulus.

Using the 1D scale to evaluate morality

With this theoretical framework, Thurstone operationalized it to learn scales for a variety of psychological phenomena but most famously known for evaluating morality of various crimes [6].

In that study, Thurstone used the scale to evaluate the public's perception of morality for 19 crimes to build a scale that could capture the moral judgments of individuals. He asked a total of 266 participants to compare every pair of crimes (171 in total) and used the entire aggregated data to learn the scale.

The resulting scale clearly had semantic sense but must be viewed from the late 1920s moral perspective.

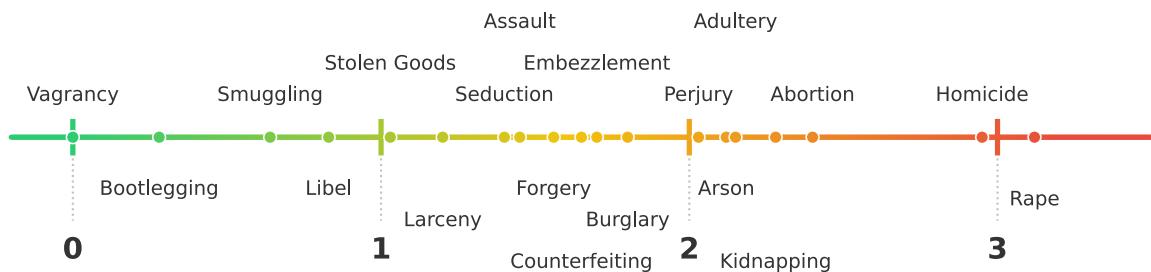


Figure 4: Thurstone's Crowdsourced 1D Scale of Seriousness of Various Crimes (1927) [6]

Serious crimes like Rape and Homicide are viewed as most immoral while minor crimes like Bootlegging (selling alcohol - this was Prohibition mind you) and Vagrancy are viewed as least immoral. One could imagine that today, this scale could look very different.

This method was revolutionary for its time being able to quantitatively model something as complex as morality without making assumptions about which crime was worse than another. Additionally, It was used for a whole other set of scales to evaluate other psychological phenomena like handwriting and personality traits [7].

Limitations of the 1D scale

However, there were definitely a few issues with this method.

1. **The scale is unidimensional:** Take the following example: it is not clear how one would compare the chocolate cake to chocolate croissant or the pound cake in terms of taste? Do you compare on the texture (cakeness) or to the taste (chocolatiness)? Thurstone's 1D method assumes that if people are confused between two items it is because they are identical on the single dimension being measured. Clearly this is not the case for this example. You lose out on this nuance in the 1D scale.
2. **Intransitivity:** Thurstone's method also assumes that if $A > B$ and $B > C$ then A must be greater than C . For example, if one were to evaluate on the chocolatiness axis then the order would likely be Chocolate Cake > Chocolate Croissant > Pound Cake. However, if one were to evaluate on the cakeness axis then the order would be Pound Cake > Chocolate Cake > Chocolate Croissant. A participant may invariably choose to evaluate on one axis for one query (Chocolate

cake vs Chocolate croissant on chocolatiness) but not for another (Pound cake vs Chocolate cake on cakeness). This subtlety is lost in the 1D scale.

3. **Lack of Scalability:** In terms of sample complexity, this is $O(N^2)$ where N is the number of stimuli. This scales quite poorly with the number of stimuli and quickly becomes infeasible in human data collection. Thurstone managed to get 266 participants to do 171 comparisons each. That in itself is a feat especially for 1927! For 25 stimuli this would 625 comparisons and for 50 one would need to collect 2500 comparisons! Good luck with that.

Coming Up Next....

Thurstone gave us the psychological ruler, but he kept us trapped on a single line. In the next episode, we'll see how researchers finally "broke out" of the first dimension and developed methods to build multidimensional psychological spaces. Episode 2 of From Scales to Spaces is titled "From Ruler to Map" and will be released next.

Stay tuned!

References

Bibliography

- [1] J. Michell, *Measurement in psychology: A critical history of a methodological concept*, vol. 53. Cambridge University Press, 1999.
- [2] A. Ferguson *et al.*, "Quantitative Estimates of Sensory Events: Final Report of the Committee appointed to consider and report upon the possibility of Quantitative Estimates of Sensory Events," *Advancement of Science*, vol. 2, pp. 331–349, 1940.
- [3] L. L. Thurstone, "A law of comparative judgment," *Scaling*. Routledge, pp. 81–92, 2017.
- [4] N. Stewart, G. D. Brown, and N. Chater, "Absolute identification by relative judgment.," *Psychological review*, vol. 112, no. 4, p. 881, 2005.
- [5] E. W. Weisstein, "Normal Difference Distribution." [Online]. Available: <https://mathworld.wolfram.com/NormalDifferenceDistribution.html>
- [6] L. L. Thurstone, "The method of paired comparisons for social values.," *The Journal of Abnormal and Social Psychology*, vol. 21, no. 4, p. 384, 1927.
- [7] L. L. Thurstone, "The measurement of social attitudes.," *The journal of abnormal and social psychology*, vol. 26, no. 3, p. 249, 1931.